

DummyRegressor

```
class sklearn.dummy.DummyRegressor(*, strategy='mean', constant=None,
quantile=None)
```

[\[source\]](#)

Regressor that makes predictions using simple rules.

This regressor is useful as a simple baseline to compare with other (real) regressors. Do not use it for real problems.

Read more in the [User Guide](#).

! Added in version 0.13.

Parameters:

strategy : {"mean", "median", "quantile", "constant"}, default="mean"

Strategy to use to generate predictions.

- "mean": always predicts the mean of the training set
- "median": always predicts the median of the training set
- "quantile": always predicts a specified quantile of the training set, provided with the quantile parameter.
- "constant": always predicts a constant value that is provided by the user.

constant : int or float or array-like of shape (n_outputs,), default=None

The explicit constant as predicted by the "constant" strategy. This parameter is useful only for the "constant" strategy.

quantile : float in [0.0, 1.0], default=None

The quantile to predict using the "quantile" strategy. A quantile of 0.5 corresponds to the median, while 0.0 to the minimum and 1.0 to the maximum.

Attributes:

constant_ : ndarray of shape (1, n_outputs)

Mean or median or quantile of the training targets or constant value given by the user.

n_features_in_ : int

Number of features seen during [fit](#).

feature_names_in_ : *ndarray of shape (n_features_in_,)*

Names of features seen during [fit](#). Defined only when `X` has feature names that are all strings.

n_outputs : *int*

Number of outputs.

➡ See also

[DummyClassifier](#)

Classifier that makes predictions using simple rules.

Examples

```
>>> import numpy as np
>>> from sklearn.dummy import DummyRegressor
>>> X = np.array([1.0, 2.0, 3.0, 4.0])
>>> y = np.array([2.0, 3.0, 5.0, 10.0])
>>> dummy_regr = DummyRegressor(strategy="mean")
>>> dummy_regr.fit(X, y)
DummyRegressor()
>>> dummy_regr.predict(X)
array([5., 5., 5., 5.])
>>> dummy_regr.score(X, y)
0.0
```

fit(*X*, *y*, *sample_weight=None*)

[\[source\]](#)

Fit the baseline regressor.

Parameters:

X : *array-like of shape (n_samples, n_features)*

Training data.

y : *array-like of shape (n_samples,) or (n_samples, n_outputs)*

Target values.

sample_weight : *array-like of shape (n_samples,), default=None*

Sample weights.

Returns:

self : object

Fitted estimator.

get_metadata_routing()

[\[source\]](#)

Get metadata routing of this object.

Please check [User Guide](#) on how the routing mechanism works.

Returns:

routing : MetadataRequest

A [MetadataRequest](#) encapsulating routing information.

get_params(deep=True)

[\[source\]](#)

Get parameters for this estimator.

Parameters:

deep : bool, default=True

If True, will return the parameters for this estimator and contained subobjects that are estimators.

Returns:

params : dict

Parameter names mapped to their values.

predict(X, return_std=False)

[\[source\]](#)

Perform classification on test vectors X.

Parameters:

X : array-like of shape (n_samples, n_features)

Test data.

return_std : bool, default=False

Whether to return the standard deviation of posterior prediction. All zeros in this case.

! Added in version 0.20.

Returns:

y : *array-like of shape (n_samples,) or (n_samples, n_outputs)*

Predicted target values for X.

y_std : *array-like of shape (n_samples,) or (n_samples, n_outputs)*

Standard deviation of predictive distribution of query points.

score(X, y, sample_weight=None)

[\[source\]](#)

Return the coefficient of determination R^2 of the prediction.

The coefficient R^2 is defined as $(1 - u/v)$, where u is the residual sum of squares $((y_{\text{true}} - y_{\text{pred}}) ** 2).sum()$ and v is the total sum of squares $((y_{\text{true}} - y_{\text{true}.mean()}) ** 2).sum()$. The best possible score is 1.0 and it can be negative (because the model can be arbitrarily worse). A constant model that always predicts the expected value of y, disregarding the input features, would get a R^2 score of 0.0.

Parameters:

X : *None or array-like of shape (n_samples, n_features)*

Test samples. Passing None as test samples gives the same result as passing real test samples, since `DummyRegressor` operates independently of the sampled observations.

y : *array-like of shape (n_samples,) or (n_samples, n_outputs)*

True values for X.

sample_weight : *array-like of shape (n_samples,), default=None*

Sample weights.

Returns:

score : *float*

R^2 of `self.predict(X)` w.r.t. y.

set_fit_request(*, sample_weight: [bool](#) | [None](#) | [str](#) = '\$UNCHANGED\$') →

[DummyRegressor](#)

[\[source\]](#)

Request metadata passed to the `fit` method.

Note that this method is only relevant if `enable_metadata_routing=True` (see [sklearn.set_config](#)). Please see [User Guide](#) on how the routing mechanism works.

The options for each parameter are:

- `True` : metadata is requested, and passed to `fit` if provided. The request is ignored if metadata is not provided.
- `False` : metadata is not requested and the meta-estimator will not pass it to `fit`.
- `None` : metadata is not requested, and the meta-estimator will raise an error if the user provides it.
- `str` : metadata should be passed to the meta-estimator with this given alias instead of the original name.

The default (`sklearn.utils.metadata_routing.UNCHANGED`) retains the existing request. This allows you to change the request for some parameters and not others.

! Added in version 1.3.

Note

This method is only relevant if this estimator is used as a sub-estimator of a meta-estimator, e.g. used inside a [Pipeline](#). Otherwise it has no effect.

Parameters:

sample_weight : *str, True, False, or None,*
default=sklearn.utils.metadata_routing.UNCHANGED

Metadata routing for `sample_weight` parameter in `fit`.

Returns:

self : *object*

The updated object.

set_params(***params*)

[\[source\]](#)

Set the parameters of this estimator.

The method works on simple estimators as well as on nested objects (such as [Pipeline](#)). The latter have parameters of the form `<component>__<parameter>` so that it's possible to update each component of a nested object.

Parameters:

****params : dict**

Estimator parameters.

Returns:

self : estimator instance

Estimator instance.

set_predict_request(*, return_std: [bool](#) / [None](#) / [str](#) = '\$UNCHANGED\$') →

[DummyRegressor](#)

[\[source\]](#)

Request metadata passed to the `predict` method.

Note that this method is only relevant if `enable_metadata_routing=True` (see [sklearn.set_config](#)). Please see [User Guide](#) on how the routing mechanism works.

The options for each parameter are:

- `True`: metadata is requested, and passed to `predict` if provided. The request is ignored if metadata is not provided.
- `False`: metadata is not requested and the meta-estimator will not pass it to `predict`.
- `None`: metadata is not requested, and the meta-estimator will raise an error if the user provides it.
- `str`: metadata should be passed to the meta-estimator with this given alias instead of the original name.

The default (`sklearn.utils.metadata_routing.UNCHANGED`) retains the existing request. This allows you to change the request for some parameters and not others.

! Added in version 1.3.

Note

This method is only relevant if this estimator is used as a sub-estimator of a meta-estimator, e.g. used inside a [Pipeline](#). Otherwise it has no effect.

Parameters:

return_std : str, True, False, or None,

default=sklearn.utils.metadata_routing.UNCHANGED

Metadata routing for `return_std` parameter in `predict`.

Returns:

self : *object*

The updated object.

`set_score_request(*, sample_weight: bool | None | str = '$UNCHANGED$') → DummyRegressor` [\[source\]](#)

Request metadata passed to the `score` method.

Note that this method is only relevant if `enable_metadata_routing=True` (see [sklearn.set_config](#)). Please see [User Guide](#) on how the routing mechanism works.

The options for each parameter are:

- `True`: metadata is requested, and passed to `score` if provided. The request is ignored if metadata is not provided.
- `False`: metadata is not requested and the meta-estimator will not pass it to `score`.
- `None`: metadata is not requested, and the meta-estimator will raise an error if the user provides it.
- `str`: metadata should be passed to the meta-estimator with this given alias instead of the original name.

The default (`sklearn.utils.metadata_routing.UNCHANGED`) retains the existing request. This allows you to change the request for some parameters and not others.

! Added in version 1.3.

Note

This method is only relevant if this estimator is used as a sub-estimator of a meta-estimator, e.g. used inside a [Pipeline](#). Otherwise it has no effect.

Parameters:

sample_weight : *str, True, False, or None,*


default=sklearn.utils.metadata_routing.UNCHANGED

Metadata routing for `sample_weight` parameter in `score`.


Returns:**self** : *object*

The updated object.

Gallery examples



Poisson regression
and non-normal
loss



Tweedie regression
on insurance claims

[Previous](#)[Next](#)

© Copyright 2007 - 2025, scikit-learn developers (BSD License).