KNeighborsClassifier

```
class sklearn.neighbors.KNeighborsClassifier(n_neighbors=5, *,
weights='uniform', algorithm='auto', leaf_size=30, p=2, metric='minkowski',
metric_params=None, n_jobs=None)
[source]
```

Classifier implementing the k-nearest neighbors vote.

Read more in the User Guide.

Parameters:

n_neighbors : *int, default=5*

Number of neighbors to use by default for kneighbors queries.

weights: {'uniform', 'distance'}, callable or None, default='uniform'

Weight function used in prediction. Possible values:

- 'uniform': uniform weights. All points in each neighborhood are weighted equally.
- 'distance': weight points by the inverse of their distance. in this case, closer neighbors of a query point will have a greater influence than neighbors which are further away.
- [callable]: a user-defined function which accepts an array of distances, and returns an array of the same shape containing the weights.

Refer to the example entitled <u>Nearest Neighbors Classification</u> showing the impact of the weights parameter on the decision boundary.

algorithm : {'auto', 'ball_tree', 'kd_tree', 'brute'}, default='auto'

Algorithm used to compute the nearest neighbors:

- 'ball_tree' will use **BallTree**
- 'kd tree' will use **KDTree**
- 'brute' will use a brute-force search.
- 'auto' will attempt to decide the most appropriate algorithm based on the values passed to **fit** method.

Note: fitting on sparse input will override the setting of this parameter, using brute force.

leaf_size : *int, default=30*

Leaf size passed to BallTree or KDTree. This can affect the speed of the construction and query, as well as the memory required to store the tree. The optimal value depends on the nature of the problem.

p: float, default=2

Power parameter for the Minkowski metric. When p = 1, this is equivalent to using manhattan_distance (I1), and euclidean_distance (I2) for p = 2. For arbitrary p, minkowski_distance (I_p) is used. This parameter is expected to be positive.

metric: str or callable, default='minkowski'

Metric to use for distance computation. Default is "minkowski", which results in the standard Euclidean distance when p = 2. See the documentation of <u>scipy.spatial.distance</u> and the metrics listed in <u>distance_metrics</u> for valid metric values.

If metric is "precomputed", X is assumed to be a distance matrix and must be square during fit. X may be a <u>sparse graph</u>, in which case only "nonzero" elements may be considered neighbors.

If metric is a callable function, it takes two arrays representing 1D vectors as inputs and must return one value indicating the distance between those vectors. This works for Scipy's metrics, but is less efficient than passing the metric name as a string.

metric_params : dict, default=None

Additional keyword arguments for the metric function.

n_jobs : int, default=None

The number of parallel jobs to run for neighbors search. None means 1 unless in a joblib.parallel_backend context. -1 means using all processors. See Glossary for more details. Doesn't affect fit method.

Attributes:

classes_ : array of shape (n_classes,)

Class labels known to the classifier

effective_metric_: str or callble

The distance metric used. It will be same as the metric parameter or a synonym of it, e.g. 'euclidean' if the metric parameter set to 'minkowski' and p parameter set to 2.

effective_metric_params_: dict

Additional keyword arguments for the metric function. For most metrics will be same with metric_params parameter, but may also contain the p parameter value if the

effective_metric_ attribute is set to 'minkowski'.

n_features_in_ : int

Number of features seen during fit.

Added in version 0.24.

feature_names_in_ : ndarray of shape (n features in ,)

Names of features seen during fit. Defined only when X has feature names that are all strings.



Added in version 1.0.

n_samples_fit_: int

Number of samples in the fitted data.

outputs_2d_ : bool

False when y's shape is (n_samples,) or (n_samples, 1) during fit otherwise True.

See also

RadiusNeighborsClassifier

Classifier based on neighbors within a fixed radius.

KNeighborsRegressor

Regression based on k-nearest neighbors.

RadiusNeighborsRegressor

Regression based on neighbors within a fixed radius.

NearestNeighbors

Unsupervised learner for implementing neighbor searches.

Notes

See Nearest Neighbors in the online documentation for a discussion of the choice of algorithm and leaf size.



Regarding the Nearest Neighbors algorithms, if it is found that two neighbors, neighbor k+1 and k, have identical distances but different labels, the results will depend on the ordering of the training data.

https://en.wikipedia.org/wiki/K-nearest_neighbor_algorithm

Examples

```
>>> X = [[0], [1], [2], [3]]
>>> y = [0, 0, 1, 1]
>>> from sklearn.neighbors import KNeighborsClassifier
>>> neigh = KNeighborsClassifier(n_neighbors=3)
>>> neigh.fit(X, y)
KNeighborsClassifier(...)
>>> print(neigh.predict([[1.1]]))
[0]
>>> print(neigh.predict_proba([[0.9]]))
[[0.666... 0.333...]]
```

fit(X, y)
[source]

Fit the k-nearest neighbors classifier from the training dataset.

Parameters:

X: {array-like, sparse matrix} of shape (n_samples, n_features) or (n_samples, n_samples) if metric='precomputed'

Training data.

y: {array-like, sparse matrix} of shape (n_samples,) or (n_samples, n_outputs)

Target values.

Returns:

self: KNeighborsClassifier

The fitted k-nearest neighbors classifier.

get_metadata_routing()

[source]

Get metadata routing of this object.

Please check User Guide on how the routing mechanism works.

Returns:

routing: MetadataRequest

A <u>MetadataRequest</u> encapsulating routing information.

get_params(deep=True)

[source]

Get parameters for this estimator.

Parameters:

deep: bool, default=True

If True, will return the parameters for this estimator and contained subobjects that are estimators.

Returns:

params : dict

Parameter names mapped to their values.

kneighbors(X=None, n_neighbors=None, return_distance=True)

[source]

Find the K-neighbors of a point.

Returns indices of and distances to the neighbors of each point.

Parameters:

X: {array-like, sparse matrix}, shape (n_queries, n_features), or (n_queries, n_indexed) if metric == 'precomputed', default=None

The query point or points. If not provided, neighbors of each indexed point are returned. In this case, the query point is not considered its own neighbor.

n_neighbors: int, default=None

Number of neighbors required for each sample. The default is the value passed to the constructor.

return distance : bool, default=True

Whether or not to return the distances.

Returns:

neigh_dist: ndarray of shape (n_queries, n_neighbors)

Array representing the lengths to points, only present if return_distance=True.

neigh_ind: ndarray of shape (n_queries, n_neighbors)

Indices of the nearest points in the population matrix.

Examples

In the following example, we construct a NearestNeighbors class from an array representing our data set and ask who's the closest point to [1,1,1]

```
>>> samples = [[0., 0., 0.], [0., .5, 0.], [1., 1., .5]]
>>> from sklearn.neighbors import NearestNeighbors
>>> neigh = NearestNeighbors(n_neighbors=1)
>>> neigh.fit(samples)
NearestNeighbors(n_neighbors=1)
>>> print(neigh.kneighbors([[1., 1., 1.]]))
(array([[0.5]]), array([[2]]))
```

As you can see, it returns [[0.5]], and [[2]], which means that the element is at distance 0.5 and is the third element of samples (indexes start at 0). You can also guery for multiple points:

```
kneighbors_graph(X=None, n_neighbors=None, mode='connectivity')
[source]
```

Compute the (weighted) graph of k-Neighbors for points in X.

Parameters:

X: {array-like, sparse matrix} of shape (n_queries, n_features), or (n_queries, n_indexed) if metric == 'precomputed', default=None

The query point or points. If not provided, neighbors of each indexed point are returned. In this case, the query point is not considered its own neighbor. For metric='precomputed' the shape should be (n_queries, n_indexed). Otherwise the shape should be (n_queries, n_features).

```
n_neighbors: int, default=None
```

Number of neighbors for each sample. The default is the value passed to the constructor.

mode: {'connectivity', 'distance'}, default='connectivity'

Type of returned matrix: 'connectivity' will return the connectivity matrix with ones and zeros, in 'distance' the edges are distances between points, type of distance depends on the selected metric parameter in NearestNeighbors class.

Returns:

A: sparse-matrix of shape (n_queries, n_samples_fit)

n_samples_fit is the number of samples in the fitted data. A[i, j] gives the weight of the edge connecting i to j. The matrix is of CSR format.

See also

NearestNeighbors.radius neighbors graph

Compute the (weighted) graph of Neighbors for points in X.

Examples

predict(x)
[source]

Predict the class labels for the provided data.

Parameters:

X: {array-like, sparse matrix} of shape (n_queries, n_features), or (n_queries, n_indexed) if metric == 'precomputed', or None

Test samples. If None, predictions for all indexed points are returned; in this case, points are not considered their own neighbors.

Returns:

y: ndarray of shape (n_queries,) or (n_queries, n_outputs)

Class labels for each data sample.

Return probability estimates for the test data X.

Parameters:

X: {array-like, sparse matrix} of shape (n_queries, n_features), or (n_queries, n_indexed) if metric == 'precomputed', or None

Test samples. If None, predictions for all indexed points are returned; in this case, points are not considered their own neighbors.

Returns:

p: ndarray of shape (n_queries, n_classes), or a list of n_outputs of such arrays if n_outputs > 1.

The class probabilities of the input samples. Classes are ordered by lexicographic order.

score(X, y, sample_weight=None)

[source]

Return the mean accuracy on the given test data and labels.

In multi-label classification, this is the subset accuracy which is a harsh metric since you require for each sample that each label set be correctly predicted.

Parameters:

X: array-like of shape (n_samples, n_features), or None

Test samples. If None, predictions for all indexed points are used; in this case, points are not considered their own neighbors. This means that knn.fit(X, y).score(None, y) implicitly performs a leave-one-out cross-validation procedure and is equivalent to cross_val_score(knn, X, y, cv=LeaveOneOut()) but typically much faster.

y: array-like of shape (n_samples,) or (n_samples, n_outputs)

True labels for x.

sample_weight: array-like of shape (n_samples,), default=None

Sample weights.

Returns:

score: float

Mean accuracy of self.predict(X) w.r.t. y.

set_params(**params)

[source]

Set the parameters of this estimator.

The method works on simple estimators as well as on nested objects (such as Pipeline). The latter have parameters of the form component> component so that it's possible to update each component of a nested object.

Parameters:

**params : dict

Estimator parameters.

Returns:

self: estimator instance

Estimator instance.

```
set_score_request(*, sample_weight: bool | None | str = '$UNCHANGED$') →
KNeighborsClassifier
[source]
```

Request metadata passed to the score method.

Note that this method is only relevant if enable_metadata_routing=True (see sklearn.set_config). Please see User Guide on how the routing mechanism works.

The options for each parameter are:

- True: metadata is requested, and passed to score if provided. The request is ignored if metadata is not provided.
- False: metadata is not requested and the meta-estimator will not pass it to score.
- None: metadata is not requested, and the meta-estimator will raise an error if the user provides it.

 str: metadata should be passed to the meta-estimator with this given alias instead of the original name.

The default (sklearn.utils.metadata_routing.UNCHANGED) retains the existing request. This allows you to change the request for some parameters and not others.

! Added in version 1.3.

1 Note

This method is only relevant if this estimator is used as a sub-estimator of a meta-estimator, e.g. used inside a **Pipeline**. Otherwise it has no effect.

Parameters:

sample_weight: str, True, False, or None,

 $default = sklearn.utils.metadata_routing.UNCHANGED$

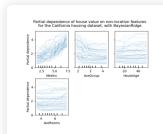
Metadata routing for sample_weight parameter in score.

Returns:

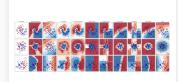
self: *object*

The updated object.

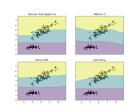
Gallery examples



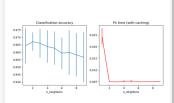
Release Highlights for scikit-learn 0.24



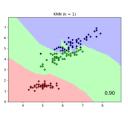
Classifier comparison



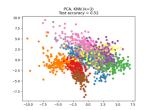
Plot the decision boundaries of a VotingClassifier



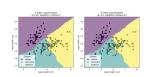
Caching nearest neighbors



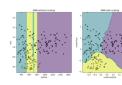
Comparing Nearest Neighbors with and without Neighborhood Components



Dimensionality Reduction with Neighborhood Components Analysis



Nearest Neighbors Classification



Importance of Feature Scaling

[©] Copyright 2007 - 2025, scikit-learn developers (BSD License).