

# recall\_score

```
sklearn.metrics.recall_score(y_true, y_pred, *, labels=None, pos_label=1,  
average='binary', sample_weight=None, zero_division='warn')
```

[\[source\]](#)

Compute the recall.

The recall is the ratio  $tp / (tp + fn)$  where `tp` is the number of true positives and `fn` the number of false negatives. The recall is intuitively the ability of the classifier to find all the positive samples.

The best value is 1 and the worst value is 0.

Support beyond term: `binary` targets is achieved by treating [multiclass](#) and [multilabel](#) data as a collection of binary problems, one for each label. For the [binary](#) case, setting `average='binary'` will return recall for `pos_label`. If `average` is not `'binary'`, `pos_label` is ignored and recall for both classes are computed then averaged or both returned (when `average=None`). Similarly, for [multiclass](#) and [multilabel](#) targets, recall for all `labels` are either returned or averaged depending on the `average` parameter. Use `labels` specify the set of labels to calculate recall for.

Read more in the [User Guide](#).

## Parameters:

**y\_true** : *1d array-like, or label indicator array / sparse matrix*

Ground truth (correct) target values.

**y\_pred** : *1d array-like, or label indicator array / sparse matrix*

Estimated targets as returned by a classifier.

**labels** : *array-like, default=None*

The set of labels to include when `average != 'binary'`, and their order if `average is None`. Labels present in the data can be excluded, for example in multiclass classification to exclude a “negative class”. Labels not present in the data can be included and will be “assigned” 0 samples. For multilabel targets, labels are column indices. By default, all labels in `y_true` and `y_pred` are used in sorted order.

❗ **Changed in version 0.17:** Parameter `labels` improved for multiclass problem.

**pos\_label** : *int, float, bool or str, default=1*

The class to report if `average='binary'` and the data is binary, otherwise this parameter is ignored. For multiclass or multilabel targets, set `labels=[pos_label]` and `average != 'binary'` to report metrics for one label only.

**`average` : {'micro', 'macro', 'samples', 'weighted', 'binary'} or None, default='binary'**

This parameter is required for multiclass/multilabel targets. If `None`, the metrics for each class are returned. Otherwise, this determines the type of averaging performed on the data:

**`'binary'` :**

Only report results for the class specified by `pos_label`. This is applicable only if targets (`y_{true,pred}`) are binary.

**`'micro'` :**

Calculate metrics globally by counting the total true positives, false negatives and false positives.

**`'macro'` :**

Calculate metrics for each label, and find their unweighted mean. This does not take label imbalance into account.

**`'weighted'` :**

Calculate metrics for each label, and find their average weighted by support (the number of true instances for each label). This alters 'macro' to account for label imbalance; it can result in an F-score that is not between precision and recall. Weighted recall is equal to accuracy.

**`'samples'` :**

Calculate metrics for each instance, and find their average (only meaningful for multilabel classification where this differs from `accuracy_score`).

**`sample_weight` : array-like of shape (n\_samples,), default=None**

Sample weights.

**`zero_division` : {"warn", 0.0, 1.0, np.nan}, default="warn"**

Sets the value to return when there is a zero division.

Notes:

- If set to "warn", this acts like 0, but a warning is also raised.
- If set to `np.nan`, such values will be excluded from the average.

**! Added in version 1.3:** `np.nan` option was added.

## Returns:

**recall** : *float (if average is not None) or array of float of shape (n\_unique\_labels,)*

Recall of the positive class in binary classification or weighted average of the recall of each class for the multiclass task.

## See also

### [precision\\_recall\\_fscore\\_support](#)

Compute precision, recall, F-measure and support for each class.

### [precision\\_score](#)

Compute the ratio  $tp / (tp + fp)$  where  $tp$  is the number of true positives and  $fp$  the number of false positives.

### [balanced\\_accuracy\\_score](#)

Compute balanced accuracy to deal with imbalanced datasets.

### [multilabel\\_confusion\\_matrix](#)

Compute a confusion matrix for each class or sample.

### [PrecisionRecallDisplay.from\\_estimator](#)

Plot precision-recall curve given an estimator and some data.

### [PrecisionRecallDisplay.from\\_predictions](#)

Plot precision-recall curve given binary class predictions.

## Notes

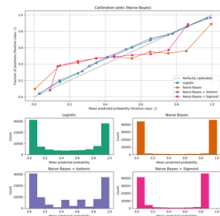
When  $true\ positive + false\ negative == 0$ , recall returns 0 and raises `UndefinedMetricWarning`. This behavior can be modified with `zero_division`.

## Examples

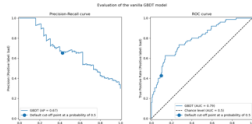
```
>>> import numpy as np
>>> from sklearn.metrics import recall_score
>>> y_true = [0, 1, 2, 0, 1, 2]
>>> y_pred = [0, 2, 1, 0, 0, 1]
>>> recall_score(y_true, y_pred, average='macro')
0.33...
>>> recall_score(y_true, y_pred, average='micro')
0.33...
>>> recall_score(y_true, y_pred, average='weighted')
0.33...
>>> recall_score(y_true, y_pred, average=None)
array([1., 0., 0.])
>>> y_true = [0, 0, 0, 0, 0, 0]
>>> recall_score(y_true, y_pred, average=None)
array([0.5, 0. , 0. ])
>>> recall_score(y_true, y_pred, average=None, zero_division=1)
array([0.5, 1. , 1. ])
>>> recall_score(y_true, y_pred, average=None, zero_division=np.nan)
array([0.5, nan, nan])
```

```
>>> # multilabel classification
>>> y_true = [[0, 0, 0], [1, 1, 1], [0, 1, 1]]
>>> y_pred = [[0, 0, 0], [1, 1, 1], [1, 1, 0]]
>>> recall_score(y_true, y_pred, average=None)
array([1. , 1. , 0.5])
```

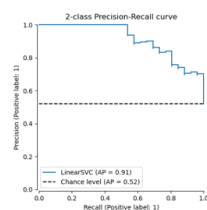
## Gallery examples



Probability



Best tuning the



Precision-Recall

© Copyright 2007 - 2025, scikit-learn developers (BSD License).